

Project 5: Recognition using Deep Networks

Team Mates: Yalala Mohit, Dhruv Kamalesh Kumar

Project Description:

The project focuses on building, training, analyzing, and modifying a deep network for a recognition task using the MNIST digit recognition dataset. The MNIST digit recognition data set will be used to train the network, which is both simple enough to be trained without a GPU and challenging enough to provide a good example of what deep networks can do.

Tasks

1. Build and train a network to recognize digits

A. Get the MNIST digit data set

The following are the first 6 mages in the MNIST digit dataset.


Ground Truth: 9



Ground Truth: 0




Ground Truth: 8




Ground Truth: 2



Ground Truth: 7



Ground Truth: 7



B. Make your network code repeatable

The code has been made repeatable by using,

```
torch.manual_seed(2502)
```

```
torch.backends.cudnn.enabled = False
```

C. Build a network model


The following is the model summary, and a visualization of the model through Netron.app

```
NeuralNetwork(  
  (conv1): Conv2d(1, 10, kernel_size=(5, 5), stride=(1, 1))  
  (conv2): Conv2d(10, 20, kernel_size=(5, 5), stride=(1, 1))  
  (conv2_drop): Dropout2d(p=0.5, inplace=False)  
  (fc1): Linear(in_features=320, out_features=50, bias=True)  
  (fc2): Linear(in_features=50, out_features=10, bias=True)  
)
```



D. Train the model

The negative log likelihood loss, is attached below. Both the train loss and test loss gradually decrease over epochs, and donot diverge. Hence, we can suggest that there is no overfitting.



E. Save the network to a file

The network is saved in various formats. Even torch script .pt format was explored and visualized.

F. Read the network and run it on the test set

The following predictions were made, when run on the test set.

Prediction: 9



Prediction: 0



Prediction: 6




Prediction: 2



Prediction: 7



Prediction: 7



Prediction: 2



Prediction: 5



Prediction: 7



G. Test the network on new inputs

The following predictions were made on our hand written digits.

Test Accuracy = 60%

Prediction: 6



Prediction: 5



Prediction: 0



Prediction: 8



Prediction: 2



Prediction: 4



Prediction: 2



Prediction: 3




Prediction: 4



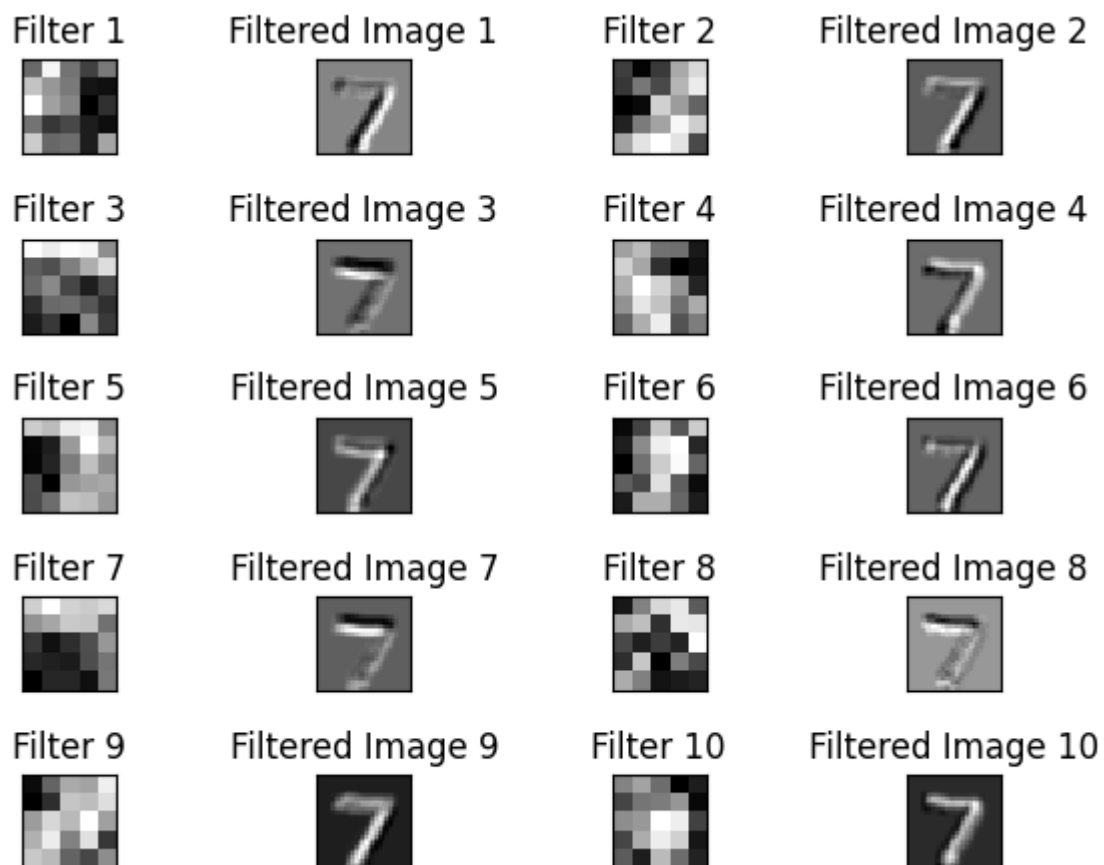
Prediction: 6



2. Examine your network
A. Analyze the first layer



B. Show the effect of the filters




The filters after being applied to the image, highlights the number in various backgrounds, extracting the features of the number. It particularly makes sense, because they're the most important features in this aspect.

3. Transfer Learning on Greek Letters

How many epochs does it take using the 27 examples in order to perfectly identify them?

The model trains and improves accuracy until 85 epochs, but is stagnated after that , even when run until 300 epochs.

The loss plot below states that overfitting has been handled. **This can be attributed to the implementation of L2 regularization during training.**



The following are the predictions made on the test set, and hand written test set.

Test set accuracy is 79% and hand written test set accuracy is 75%

Predicted: Alpha



Predicted: Alpha



Predicted: Beta



Predicted: Alpha



Predicted: Gamma



Predicted: Beta



Predicted: Alpha




Predicted: Gamma




Predicted: Alpha




Predicted Train: Alpha




Predicted Train: Alpha




Predicted Train: Alpha



Predicted Train: Gamma



Predicted Train: Gamma



4. Design your own experiment

A. Develop a plan

The plan is to evaluate the following Hyperparameters. The model will evaluate all $L \times M \times N$ combinations, which will correspond to 1024 models eventually.

```
# put all hyper params into a OrderedDict, easily expandable
params = OrderedDict([
    lr = [.01, .001],
    batch_size = [100, 1000],
    shuffle = [True],
    epochs = [5,10],
    conv_channels = [[6,12], [8,16]],
    conv_kernel_size = [3,5],
    pool_kernel_size = [2,3],
    pool_stride = [1, 2],
    dropout_rate = [0.1, 0.5],
    hidden_layers = [[120, 60], [512,256]],
    activation = [nn.ReLU(), nn.Tanh()],
])
```

B. Predict the results

The hypothesis is that the learning rate of 0.001 will work the best with 10 epochs, and batch size of 100 and hidden layer rate as [120,60], with Relu activation.


The rest of the parameters cannot be predicted.

The assumption was that after a few epochs, the model was already working very well. Hence, the model doesn't need to have extremely huge parameters.



C. Execute your plan

As expected the parameters that perform the best are in accordance with the hypothesis.

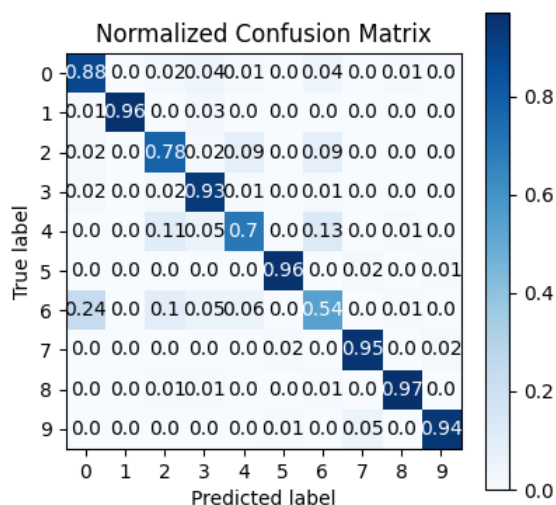
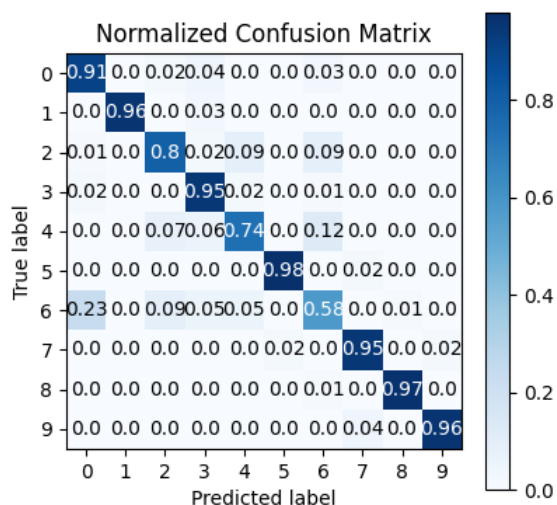
The confusion matrixes of top 5 working models have been shown below.





Train accuracy = 0.8926166666666666, Test accuracy = 0.8762




Train accuracy = 0.8439166666666666, Test accuracy = 0.8297



Train accuracy = 0.8798333333333334, Test accuracy = 0.8619



Train accuracy = 0.8669333333333333, Test accuracy = 0.8535



Train accuracy = 0.8633166666666666, Test accuracy = 0.8492

5. Extensions

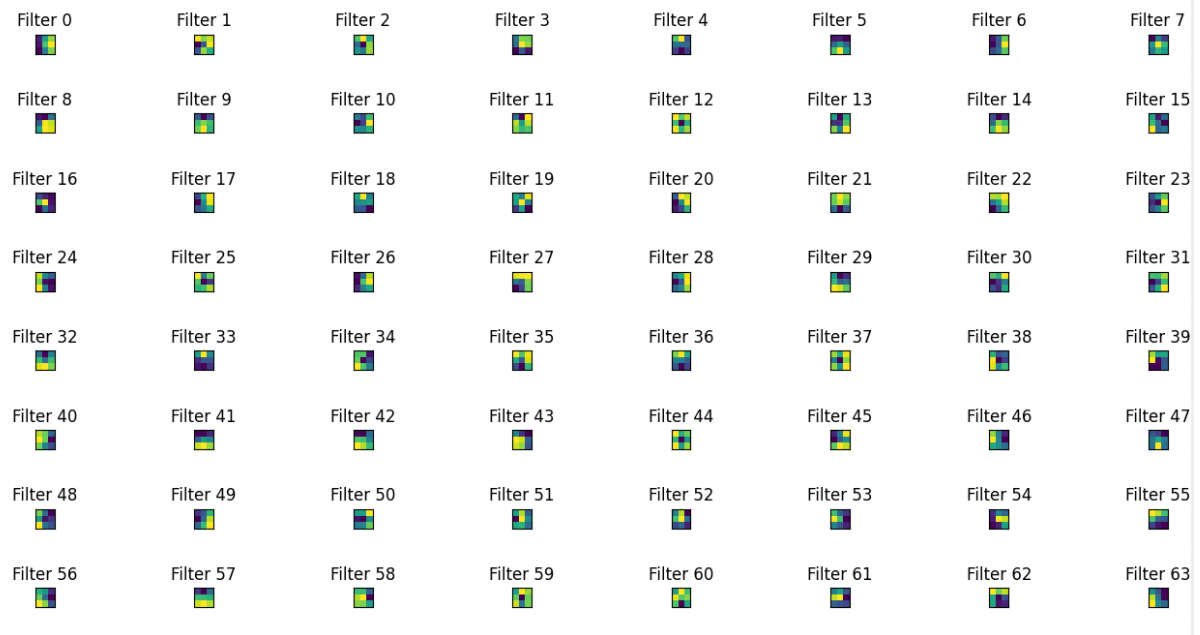
- A. There are many pre-trained networks available in the PyTorch package. Try loading one and evaluate its first couple of convolutional layers as in task 2.

Model: VGG16.

Test images – 2 images from CIFAR 10 dataset.

- Shape of layer 1 of the network - torch.Size([64, 3, 3, 3])
- Shape of first filter - torch.Size([3, 3])
- First filter – tensor([
- [-0.5537, 0.1427, 0.5290],
- [-0.5831, 0.3566, 0.7657],
- [-0.6902, -0.0480, 0.4841]], grad_fn=<SelectBackward0>)
- Shape of layer 3 of the network - torch.Size([64, 64, 3, 3])
- Shape of first filter - torch.Size([3, 3])
- First filter - tensor([
- [-0.0306, -0.0985, -0.1326],
- [0.0068, -0.0835, -0.1670],
- [0.0310, -0.0658, -0.1317]], grad_fn=<SelectBackward0>)

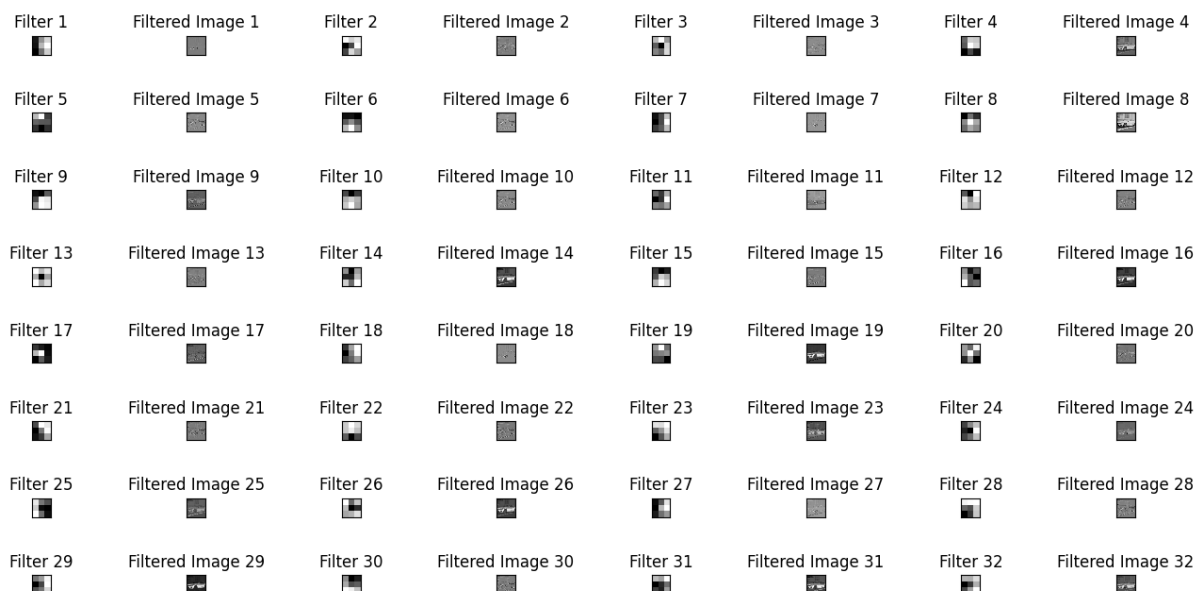
Conv 1- filters




Conv2 -Filters



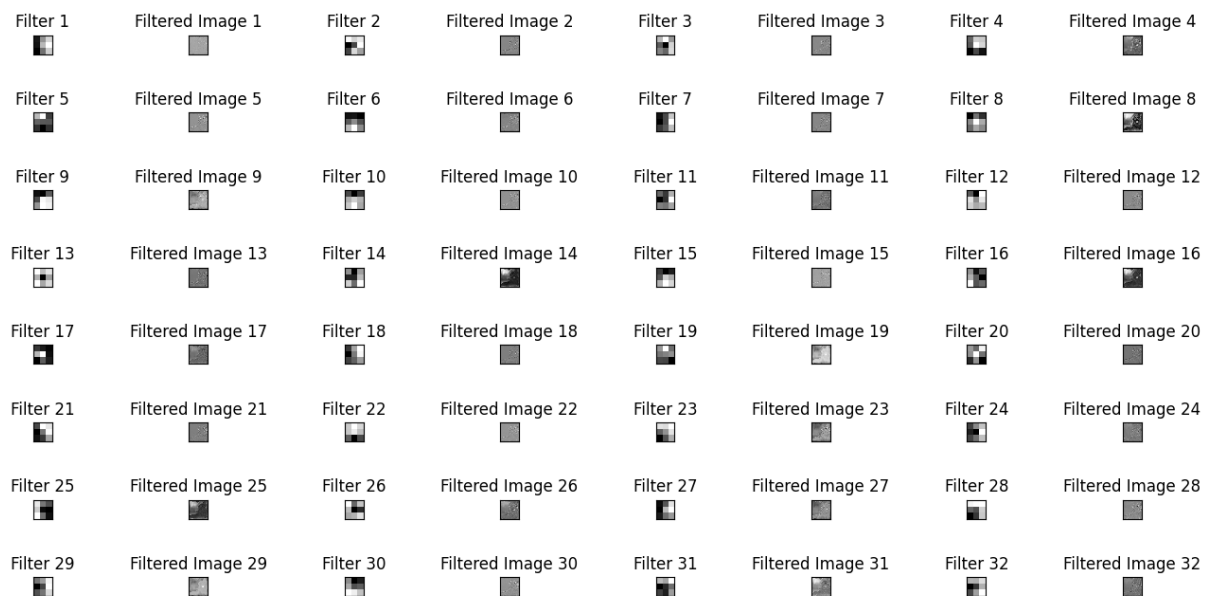
Conv1- filter applied to image – Car from CIFAR10 dataset



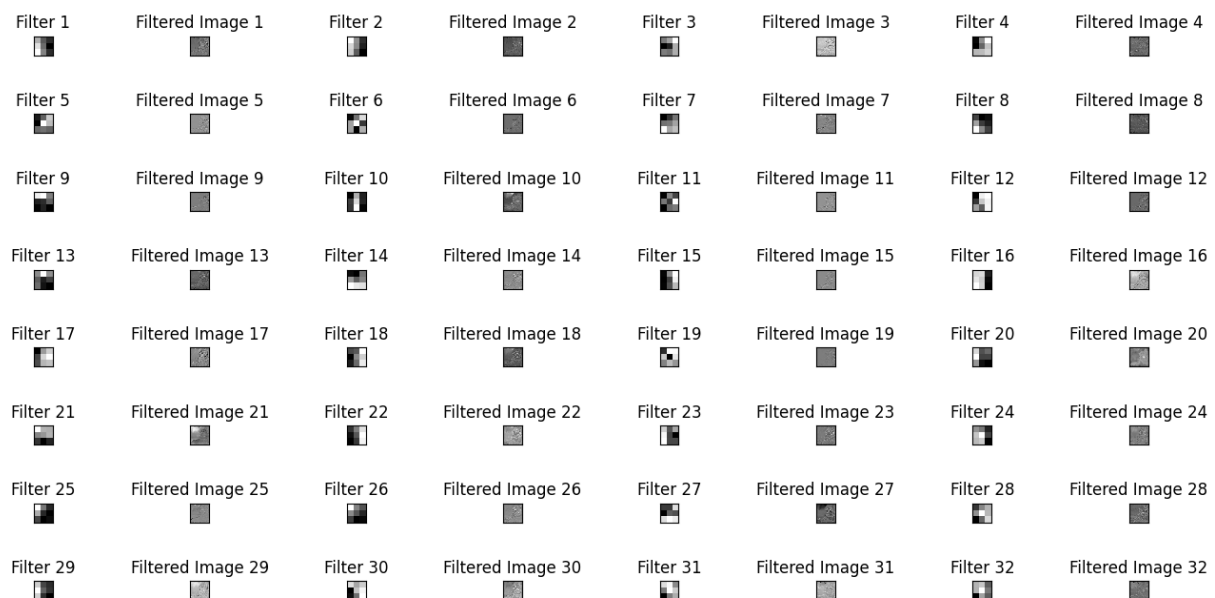
Conv2- filter applied to image – Car from CIFAR10 dataset



Conv1- filter applied to image – Dog from CIFAR10 dataset



Conv2- filter applied to image – Dog from CIFAR10 dataset



- B. Replace the first layer of the MNIST network with a filter bank of your choosing (e.g. Gabor filters) and retrain the rest of the network, holding the first layer constant. How does it do?

We applied 3 different kinds of filters,

- Gabor Filter
- Laplacian Filter
- Gaussian Filter

The Gabor Filter when applied to first conv layer and frozen performs really bad, with an accuracy of 11% which was not anticipated.

24	Gabor model training		
25	Train Epoch: 1 [0/60000 (0%)]	Loss: 2.304324	
26	Train Epoch: 1 [10000/60000 (17%)]	Loss: 2.293156	
27	Train Epoch: 1 [20000/60000 (33%)]	Loss: 2.295112	
28	Train Epoch: 1 [30000/60000 (50%)]	Loss: 2.292838	
29	Train Epoch: 1 [40000/60000 (67%)]	Loss: 2.296666	
30	Train Epoch: 1 [50000/60000 (83%)]	Loss: 2.301371	
31			
32	Test set: Avg. loss: 2.3012, Accuracy: 1135/10000 (11%)		
33			
34	Train Epoch: 2 [0/60000 (0%)]	Loss: 2.301706	
35	Train Epoch: 2 [10000/60000 (17%)]	Loss: 2.300179	
36	Train Epoch: 2 [20000/60000 (33%)]	Loss: 2.295363	
37	Train Epoch: 2 [30000/60000 (50%)]	Loss: 2.308613	
38	Train Epoch: 2 [40000/60000 (67%)]	Loss: 2.303169	
39	Train Epoch: 2 [50000/60000 (83%)]	Loss: 2.301089	
40			
41	Test set: Avg. loss: 2.3011, Accuracy: 1135/10000 (11%)		
42			
43	Train Epoch: 3 [0/60000 (0%)]	Loss: 2.298507	
44	Train Epoch: 3 [10000/60000 (17%)]	Loss: 2.313024	
45	Train Epoch: 3 [20000/60000 (33%)]	Loss: 2.304189	
46	Train Epoch: 3 [30000/60000 (50%)]	Loss: 2.291450	
47	Train Epoch: 3 [40000/60000 (67%)]	Loss: 2.296810	
48	Train Epoch: 3 [50000/60000 (83%)]	Loss: 2.298209	
49			
50	Test set: Avg. loss: 2.3011, Accuracy: 1135/10000 (11%)		
51			
52	Train Epoch: 4 [0/60000 (0%)]	Loss: 2.300630	
53	Train Epoch: 4 [10000/60000 (17%)]	Loss: 2.300376	
54	Train Epoch: 4 [20000/60000 (33%)]	Loss: 2.310190	
55	Train Epoch: 4 [30000/60000 (50%)]	Loss: 2.293983	
56	Train Epoch: 4 [40000/60000 (67%)]	Loss: 2.292970	
57	Train Epoch: 4 [50000/60000 (83%)]	Loss: 2.303795	
58			
59	Test set: Avg. loss: 2.3011, Accuracy: 1135/10000 (11%)		
60			
61	Train Epoch: 5 [0/60000 (0%)]	Loss: 2.311693	
62	Train Epoch: 5 [10000/60000 (17%)]	Loss: 2.311175	
63	Train Epoch: 5 [20000/60000 (33%)]	Loss: 2.308060	
64	Train Epoch: 5 [30000/60000 (50%)]	Loss: 2.304125	
65	Train Epoch: 5 [40000/60000 (67%)]	Loss: 2.297257	
66	Train Epoch: 5 [50000/60000 (83%)]	Loss: 2.302872	
67			
68	Test set: Avg. loss: 2.3011, Accuracy: 1135/10000 (11%)		

The Laplacian filter performs fairly better with an accuracy of 93%


Laplacian model training		
Train Epoch: 1 [0/60000 (0%)]	Loss: 2.078504	
Train Epoch: 1 [10000/60000 (17%)]	Loss: 0.943062	
Train Epoch: 1 [20000/60000 (33%)]	Loss: 1.045826	
Train Epoch: 1 [30000/60000 (50%)]	Loss: 1.013499	
Train Epoch: 1 [40000/60000 (67%)]	Loss: 0.922272	
Train Epoch: 1 [50000/60000 (83%)]	Loss: 0.686135	
Test set: Avg. loss: 0.3931, Accuracy: 8892/10000 (89%)		
Train Epoch: 2 [0/60000 (0%)]	Loss: 0.559149	
Train Epoch: 2 [10000/60000 (17%)]	Loss: 0.724622	
Train Epoch: 2 [20000/60000 (33%)]	Loss: 0.720831	
Train Epoch: 2 [30000/60000 (50%)]	Loss: 0.678157	
Train Epoch: 2 [40000/60000 (67%)]	Loss: 0.700722	
Train Epoch: 2 [50000/60000 (83%)]	Loss: 0.723783	
Test set: Avg. loss: 0.3109, Accuracy: 9097/10000 (91%)		
Train Epoch: 3 [0/60000 (0%)]	Loss: 0.697680	
Train Epoch: 3 [10000/60000 (17%)]	Loss: 0.523949	
Train Epoch: 3 [20000/60000 (33%)]	Loss: 0.569957	
Train Epoch: 3 [30000/60000 (50%)]	Loss: 0.607950	
Train Epoch: 3 [40000/60000 (67%)]	Loss: 0.568528	
Train Epoch: 3 [50000/60000 (83%)]	Loss: 0.527512	
Test set: Avg. loss: 0.2870, Accuracy: 9178/10000 (92%)		
Train Epoch: 4 [0/60000 (0%)]	Loss: 0.705846	
Train Epoch: 4 [10000/60000 (17%)]	Loss: 0.739026	
Train Epoch: 4 [20000/60000 (33%)]	Loss: 0.686989	
Train Epoch: 4 [30000/60000 (50%)]	Loss: 0.552179	
Train Epoch: 4 [40000/60000 (67%)]	Loss: 0.764357	
Train Epoch: 4 [50000/60000 (83%)]	Loss: 0.617486	
Test set: Avg. loss: 0.2487, Accuracy: 9257/10000 (93%)		

The Gaussian filter works the best with an accuracy of 95%.

Gaussian model training		
Train Epoch: 1	[0/60000 (0%)]	Loss: 1.609859
Train Epoch: 1	[10000/60000 (17%)]	Loss: 0.788865
Train Epoch: 1	[20000/60000 (33%)]	Loss: 0.501100
Train Epoch: 1	[30000/60000 (50%)]	Loss: 0.470818
Train Epoch: 1	[40000/60000 (67%)]	Loss: 0.407144
Train Epoch: 1	[50000/60000 (83%)]	Loss: 0.288379
Test set: Avg. loss: 0.1849, Accuracy: 9472/10000 (95%)		
Train Epoch: 2	[0/60000 (0%)]	Loss: 0.336808
Train Epoch: 2	[10000/60000 (17%)]	Loss: 0.433795
Train Epoch: 2	[20000/60000 (33%)]	Loss: 0.326200
Train Epoch: 2	[30000/60000 (50%)]	Loss: 0.542895
Train Epoch: 2	[40000/60000 (67%)]	Loss: 0.366523
Train Epoch: 2	[50000/60000 (83%)]	Loss: 0.311382
Test set: Avg. loss: 0.1565, Accuracy: 9529/10000 (95%)		
Train Epoch: 3	[0/60000 (0%)]	Loss: 0.565052
Train Epoch: 3	[10000/60000 (17%)]	Loss: 0.396134
Train Epoch: 3	[20000/60000 (33%)]	Loss: 0.400395
Train Epoch: 3	[30000/60000 (50%)]	Loss: 0.336715
Train Epoch: 3	[40000/60000 (67%)]	Loss: 0.351089
Train Epoch: 3	[50000/60000 (83%)]	Loss: 0.350298
Test set: Avg. loss: 0.1393, Accuracy: 9569/10000 (96%)		
Train Epoch: 4	[0/60000 (0%)]	Loss: 0.274269
Train Epoch: 4	[10000/60000 (17%)]	Loss: 0.316038
Train Epoch: 4	[20000/60000 (33%)]	Loss: 0.293957
Train Epoch: 4	[30000/60000 (50%)]	Loss: 0.385821
Train Epoch: 4	[40000/60000 (67%)]	Loss: 0.361743
Train Epoch: 4	[50000/60000 (83%)]	Loss: 0.246931
Train Epoch: 2	[10000/60000 (17%)]	Loss: 0.433795
Train Epoch: 2	[20000/60000 (33%)]	Loss: 0.326200
Train Epoch: 2	[30000/60000 (50%)]	Loss: 0.542895
Train Epoch: 2	[40000/60000 (67%)]	Loss: 0.366523
Train Epoch: 2	[50000/60000 (83%)]	Loss: 0.311382
Test set: Avg. loss: 0.1565, Accuracy: 9529/10000 (95%)		
Train Epoch: 3	[0/60000 (0%)]	Loss: 0.565052
Train Epoch: 3	[10000/60000 (17%)]	Loss: 0.396134
Train Epoch: 3	[20000/60000 (33%)]	Loss: 0.400395
Train Epoch: 3	[30000/60000 (50%)]	Loss: 0.336715
Train Epoch: 3	[40000/60000 (67%)]	Loss: 0.351089
Train Epoch: 3	[50000/60000 (83%)]	Loss: 0.350298
Test set: Avg. loss: 0.1393, Accuracy: 9569/10000 (96%)		
Train Epoch: 4	[0/60000 (0%)]	Loss: 0.274269
Train Epoch: 4	[10000/60000 (17%)]	Loss: 0.316038
Train Epoch: 4	[20000/60000 (33%)]	Loss: 0.293957
Train Epoch: 4	[30000/60000 (50%)]	Loss: 0.385821
Train Epoch: 4	[40000/60000 (67%)]	Loss: 0.361743
Train Epoch: 4	[50000/60000 (83%)]	Loss: 0.246931
Test set: Avg. loss: 0.1180, Accuracy: 9632/10000 (96%)		

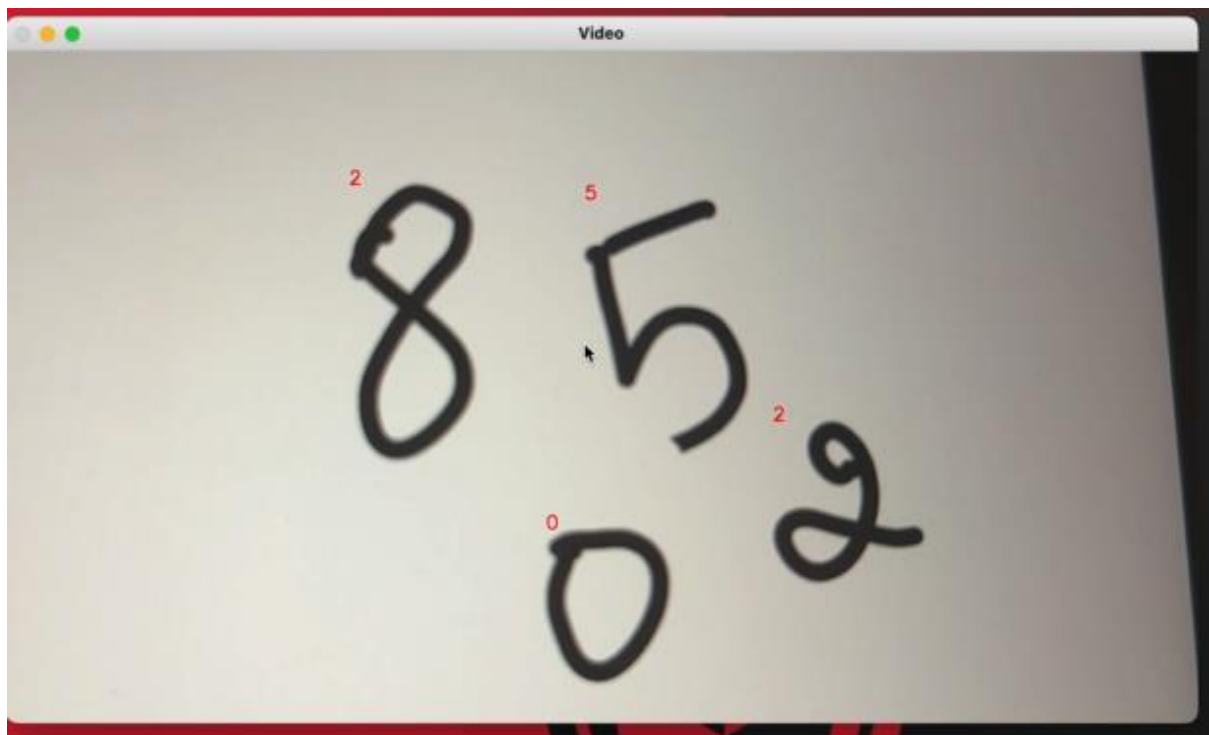
C. Build a live video digit recognition application using the trained network.

The application was developed, to recognize a single digit in the live video stream.



D. Build a live video Multi digit recognition application using the trained network.

The previous extension was again extended to recognize Multiple digits in the same frame.



E. Evaluate more dimensions on task 4 – 1024 Models.

A total of 1024 models were evaluated, with the best performing model having an accuracy of 87% on the Fashion MNIST dataset. The results for all runs with all hyperparameters can be seen in results.csv and results.json

6. Short Reflection

The project allowed us to understand all the intricacies involved in pytorch. We were both extensive users of tensorflow and were new to pytorch, but now we're able to understand and apply all techniques in Pytorch too. The library is very effective and provided great flexibility.

Additionally, it was very useful to learn hyper parameter tuning, and test out various hypothesis.

7. References

- PyTorch documentation - <https://pytorch.org/docs/stable/index.html>
- PyTorch tutorials - <https://pytorch.org/tutorials/>
- torchvision documentation - <https://pytorch.org/vision/stable/index.html>
- MNIST dataset website - <http://yann.lecun.com/exdb/mnist/>
- Matplotlib documentation - <https://matplotlib.org/stable/contents.html>
- OpenCV documentation - <https://docs.opencv.org/master/>
- NumPy documentation - <https://numpy.org/doc/stable/>
- Pandas documentation - <https://pandas.pydata.org/docs/>
- Seaborn documentation - <https://seaborn.pydata.org/documentation.html>
- SciPy documentation - <https://docs.scipy.org/doc/>
- Jupyter Notebook documentation - <https://jupyter-notebook.readthedocs.io/en/stable/>
- Google Colab documentation - <https://colab.research.google.com/notebooks/intro.ipynb>
- GitHub documentation - <https://docs.github.com/en>
- Git documentation - <https://git-scm.com/doc>